Oregon State University
CS512, Winter 2025
Group Homework: BigQuery - Bigish Data
Team Nexus
Paul J Anderson
Rachel Hughes

Date 2.15.2025

This is a Big Data assignment that served to give us experience with setting up and running Linux-based VMs from a Unix SSH terminal, loading large datasets, scrubbing our data using cloud tools, then doing initial investigation into trends and patterns and answering simple questions. For this assignment, we used Google Cloud Services (Compute Engine, Cloud Dataprep, and BigQuery) and a 10GB datafile of worldwide ICAO transmissions for a given period of time.

Question: "How many unique airplanes identified by their ICAO were in the area defined by Lat 44.497222 +/- 0.2 and Lon -123.289444 +/- 0.2. This is a rectangle around the Corvallis, OR airport?"
ANSWER: 85

1. Get Zip file onto Compute Engine Instance

```
andepaul@instance-20250215-065317:~$ lsblk
NAME    MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
sda       8:0    0   10G  0 disk
├─sda1    8:1    0  9.9G  0 part /
├─sda14   8:14   0    3M  0 part
└─sda15   8:15   0  124M  0 part /boot/efi
sdb       8:16   0  100G  0 disk /mnt/data
andepaul@instance-20250215-065317:~$ cd /mnt/data
andepaul@instance-20250215-065317:/mnt/data$ df -h /mnt/data
Filesystem      Size  Used Avail Use% Mounted on
/dev/sdb         98G   43G   51G  46% /mnt/data
andepaul@instance-20250215-065317:/mnt/data$ ls /mnt/data
apt-cache  lost+found  plane_data  plane_data.zip
andepaul@instance-20250215-065317:/mnt/data$
```

We created a VM using Google Compute Engine. Initially, we used a simple machine with only 4 GB of RAM and 10 GB of boot disk space. We purchased an external disk with 100 GB as an expansion. The initial download of the zip file took about 50 minutes and stopped near completion due to a full hard drive. We stopped the machine and remade it with an e2-standard-4 instance, 16 GB RAM, and a 100 GB SCSI disk, costing $98/month. Before downloading the zip file, we ensured we were using the 100 GB disk. We followed the documentation on file to use the necessary commands. This new machine downloaded the zip file in 17 minutes. The external disk space needed mapping, so we followed the directions provided in the documentation to mount the disk and alter our commands

to include the mount drive location (in our case, /mnt/data/). We were able to successfully unzip the JSON files and isolate them by creating another folder and moving all non-JSON files there.

2. Load JSON files into Google Cloud Storage.

We then created a bucket on Google Cloud Storage named cs512_aircraft and used cloud-init in the SSH terminal to initiate the cloud. We moved some files around and reviewed what files we had in each folder directory to ensure our JSONs were ready for upload. We uploaded the plane_data JSONs to our bucket, gs://aircraft, using the gsutil cp command. We did run into a permission issue that prevented 75 objects from being transferred. Otherwise, the JSON data is confirmed in Google Cloud Storage.

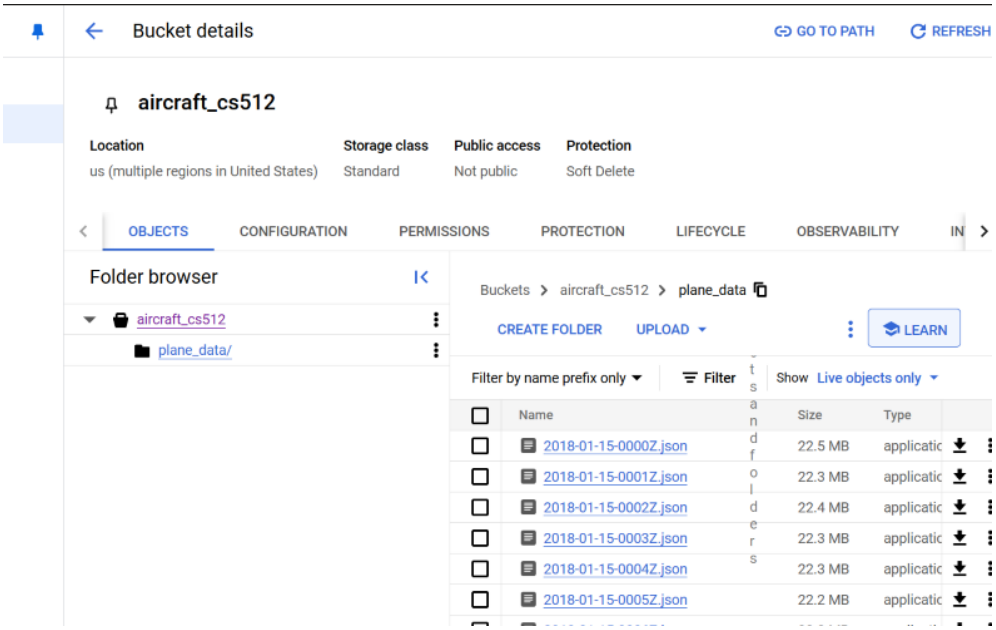

3. Load JSON files as a data set into Google Cloud Dataprep.

Next, we attempted to load the JSON files as a dataset into Dataprep. We tried this in several ways, as illustrated in our screenshots below. Although plane_data is clearly a file full of JSONs, previews of these data were not available. We created datasets in SSH gs:// and verified they were present in BigQuery, but the imported JSONs were not. Even a single JSON showed no records. We attempted to remove formatting from the JSON files and upload them again, but this returned another error. We consulted Ed, spoke with other students, messaged the instructor, and attended office hours. We were encouraged to complete the assignment without these steps. Therefore, this portion of the assignment was not completed. Significant effort was made to try and process a JSON through these steps. Assistance from team members, the professor, and the TA was sought, but a solution was not forthcoming.

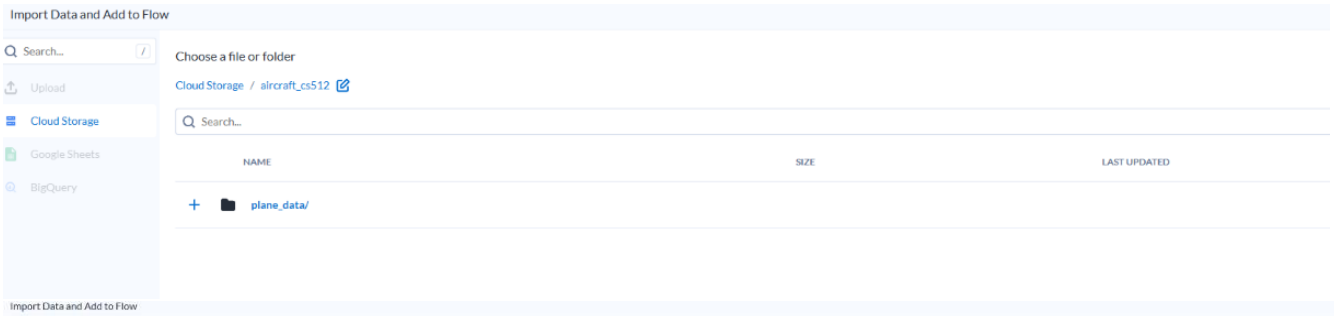
Import Data and Add to Flow
Search...
Upload
Cloud Storage
Google Sheets
BigQuery
Choose a file or folder
Cloud Storage / aircraft_cs512
Search...
NAME
SIZE
LAST UPDATED
plane_data/

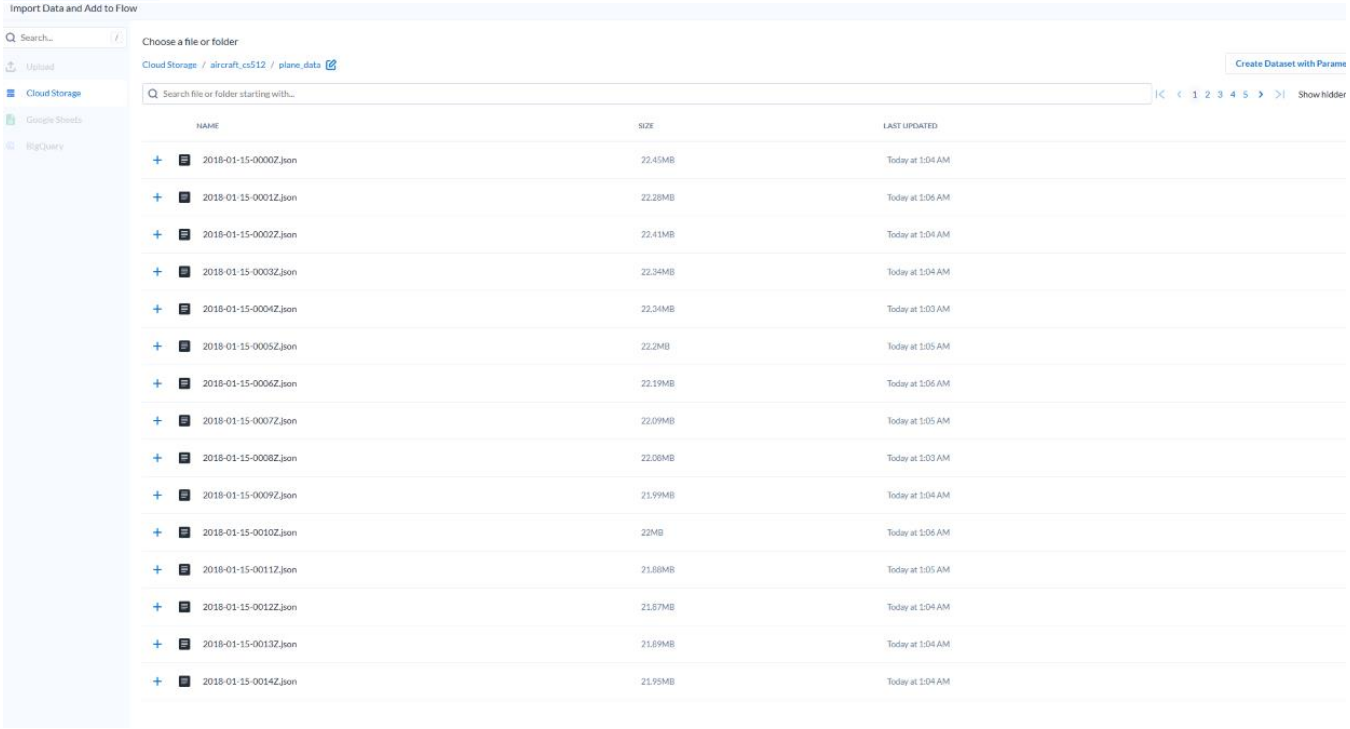
Import Data and Add to Flow
Search...
Upload
Cloud Storage
Google Sheets
BigQuery
Choose a file or folder
Cloud Storage / aircraft_cs512 / plane_data
Create Dataset with Parame
Search file or folder starting with...
1 2 3 4 5
Show hidden
NAME
SIZE
LAST UPDATED
2018-01-15-0000Z.json 22.45MB Today at 1:04 AM
2018-01-15-0001Z.json 22.28MB Today at 1:06 AM
2018-01-15-0002Z.json 22.41MB Today at 1:04 AM
2018-01-15-0003Z.json 22.34MB Today at 1:04 AM
2018-01-15-0004Z.json 22.34MB Today at 1:03 AM
2018-01-15-0005Z.json 22.2MB Today at 1:05 AM
2018-01-15-0006Z.json 22.19MB Today at 1:06 AM
2018-01-15-0007Z.json 22.09MB Today at 1:05 AM
2018-01-15-0008Z.json 22.08MB Today at 1:03 AM
2018-01-15-0009Z.json 21.99MB Today at 1:04 AM
2018-01-15-0010Z.json 22MB Today at 1:06 AM
2018-01-15-0011Z.json 21.88MB Today at 1:05 AM
2018-01-15-0012Z.json 21.87MB Today at 1:04 AM
2018-01-15-0013Z.json 21.89MB Today at 1:04 AM
2018-01-15-0014Z.json 21.95MB Today at 1:04 AM

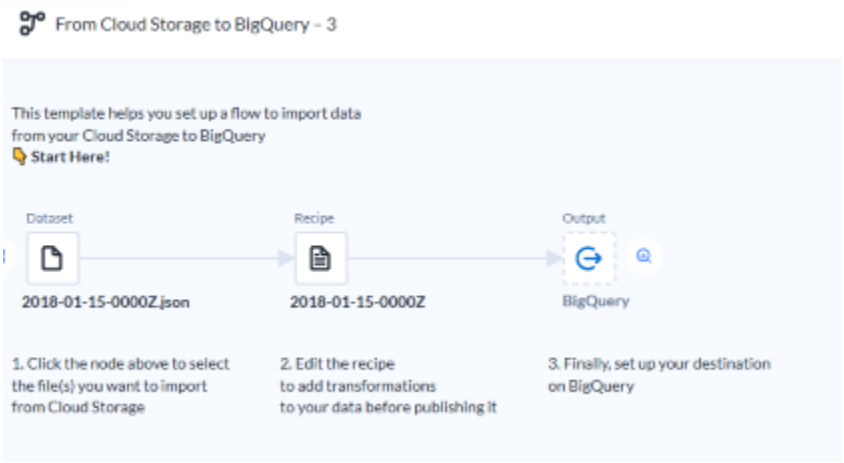
From Cloud Storage to BigQuery - 3
This template helps you set up a flow to import data from your Cloud Storage to BigQuery
Start Here!
Dataset
Recipe
Output
2018-01-15-0000Z.json
2018-01-15-0000Z
BigQuery
1. Click the node above to select the file(s) you want to import from Cloud Storage
2. Edit the recipe to add transformations to your data before publishing it
3. Finally, set up your destination on BigQuery

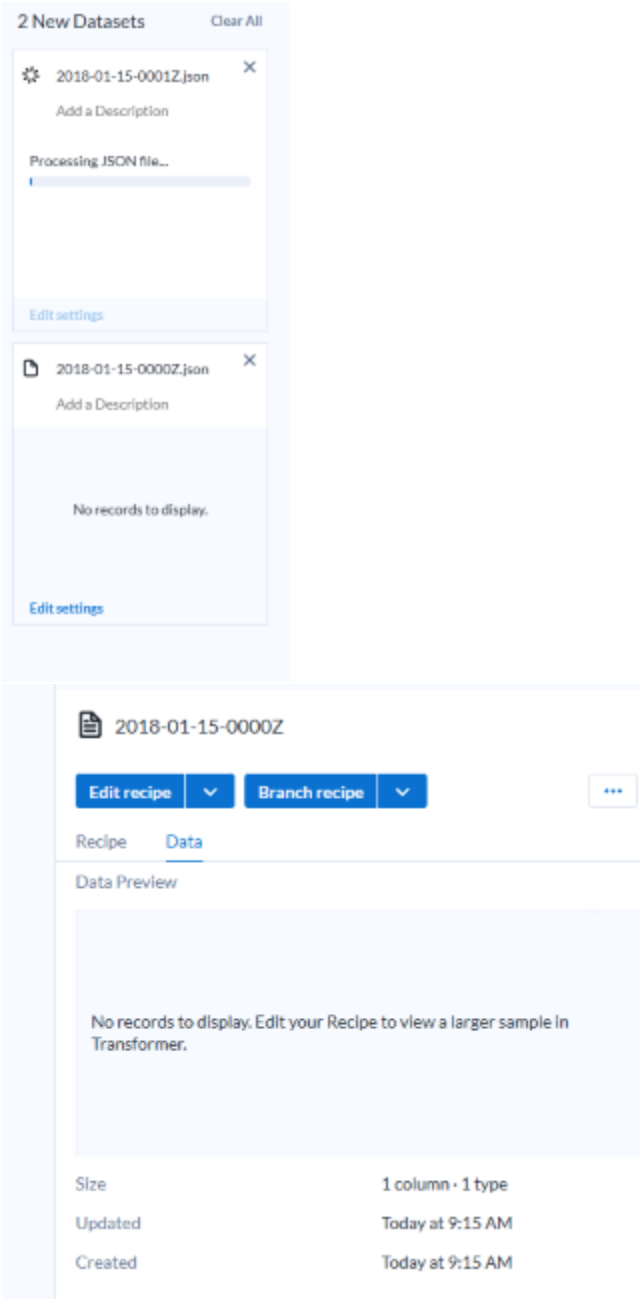
2 New Datasets
Clear All
2018-01-15-0001Z.json
Add a Description
Processing JSON file...
Edit settings
2018-01-15-0000Z.json
Add a Description
No records to display.
Edit settings
2018-01-15-0000Z
Edit recipe
Branch recipe
Recipe
Data
Data Preview
No records to display. Edit your Recipe to view a larger sample in Transformer.
Size
1 column · 1 type
Updated
Today at 9:15 AM
Created
Today at 9:15 AM

Dataprep by TRIFACTA

Cloud Dataprep by Trifacta is now **Generally Available** for all users and fully qualified for production use. This service is provided in collaboration with Google LLC. Google Cloud Platform support details can be found here.

Please review and accept the Terms of Service below to start or continue use of Cloud Dataprep.

☐ By checking this box, you agree to the Trifacta Terms of Service.

Deny & Log out | Accept

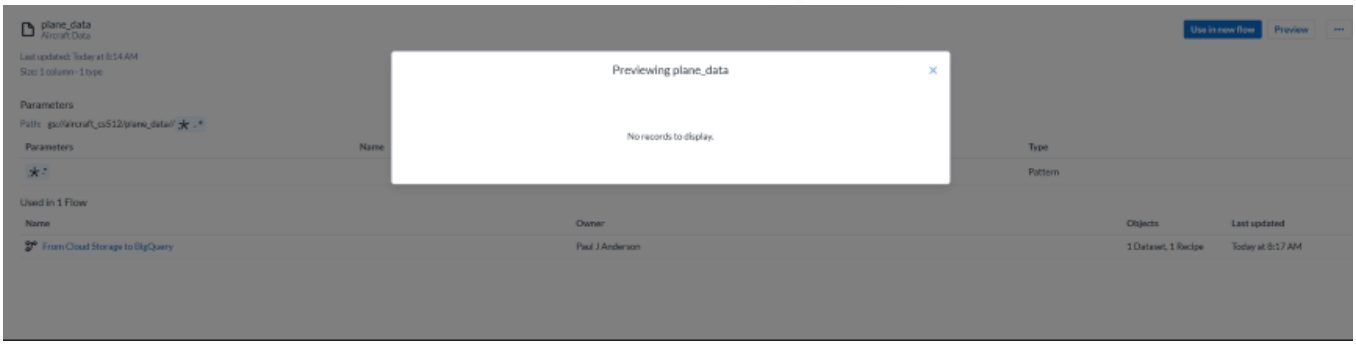

Create Output

Choose a loading option

- **Create new table** — Data will be published into a new table.
- **Replace data only (Truncate)** — Data will replace the entire content of the selected table. Table metadata will not change
- **Append to table** — Data will be inserted at the end of the selected table.
- **Drop the table** — The selected table will be deleted and replaced with a new table with the data.

Learn more about loading options

Create new table

Project: CS512 (cs512-447721)

Dataset: Choose a dataset

- aircraft_data
- test_dataset

Cancel | Save

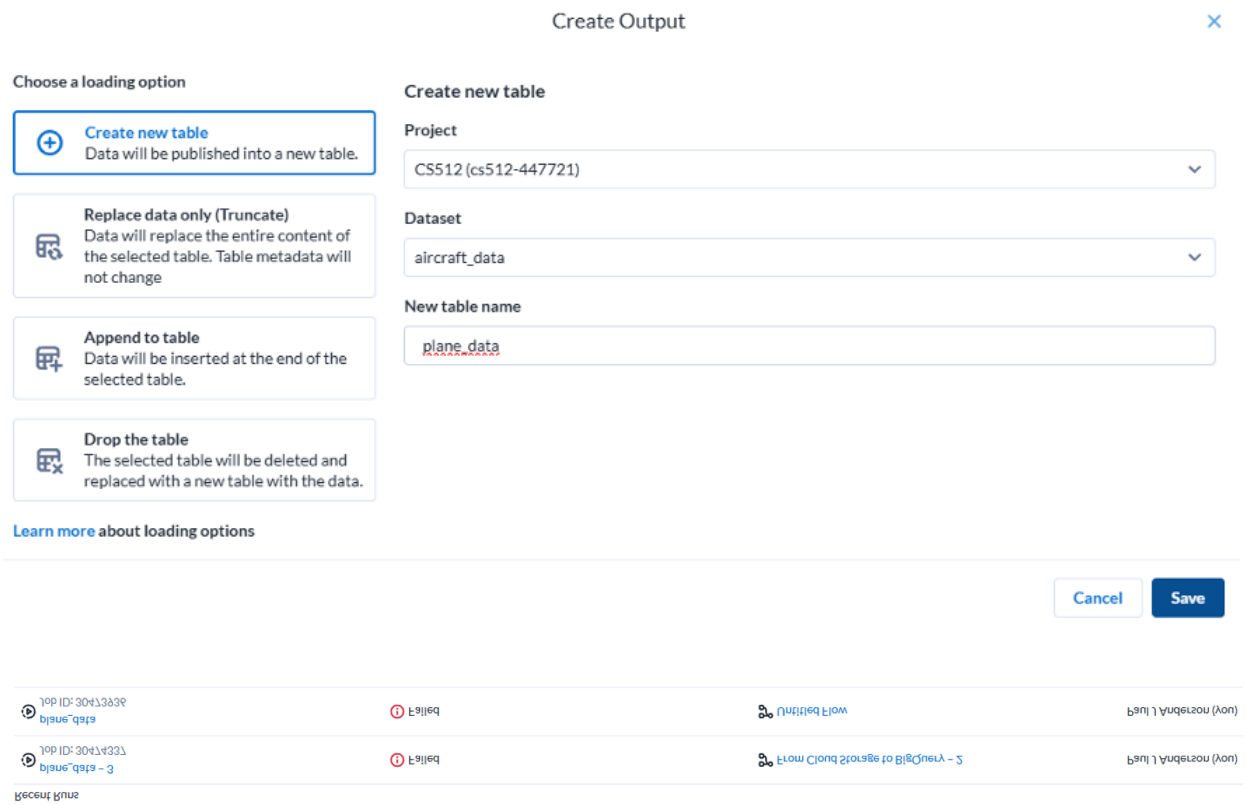

4. Parse JSON into appropriate columns in Dataprep

Like the previous step, this portion of the assignment was not completed. Significant effort was made to try and process a JSON through these steps. Assistance from team members, the professor, and the TA was sought, but a solution was not forthcoming.

5. Export Dataprep job into BigQuery

We were able to upload a pre-wrangled .CSV into BigQuery via the following steps. Create table from a GCS bucket.

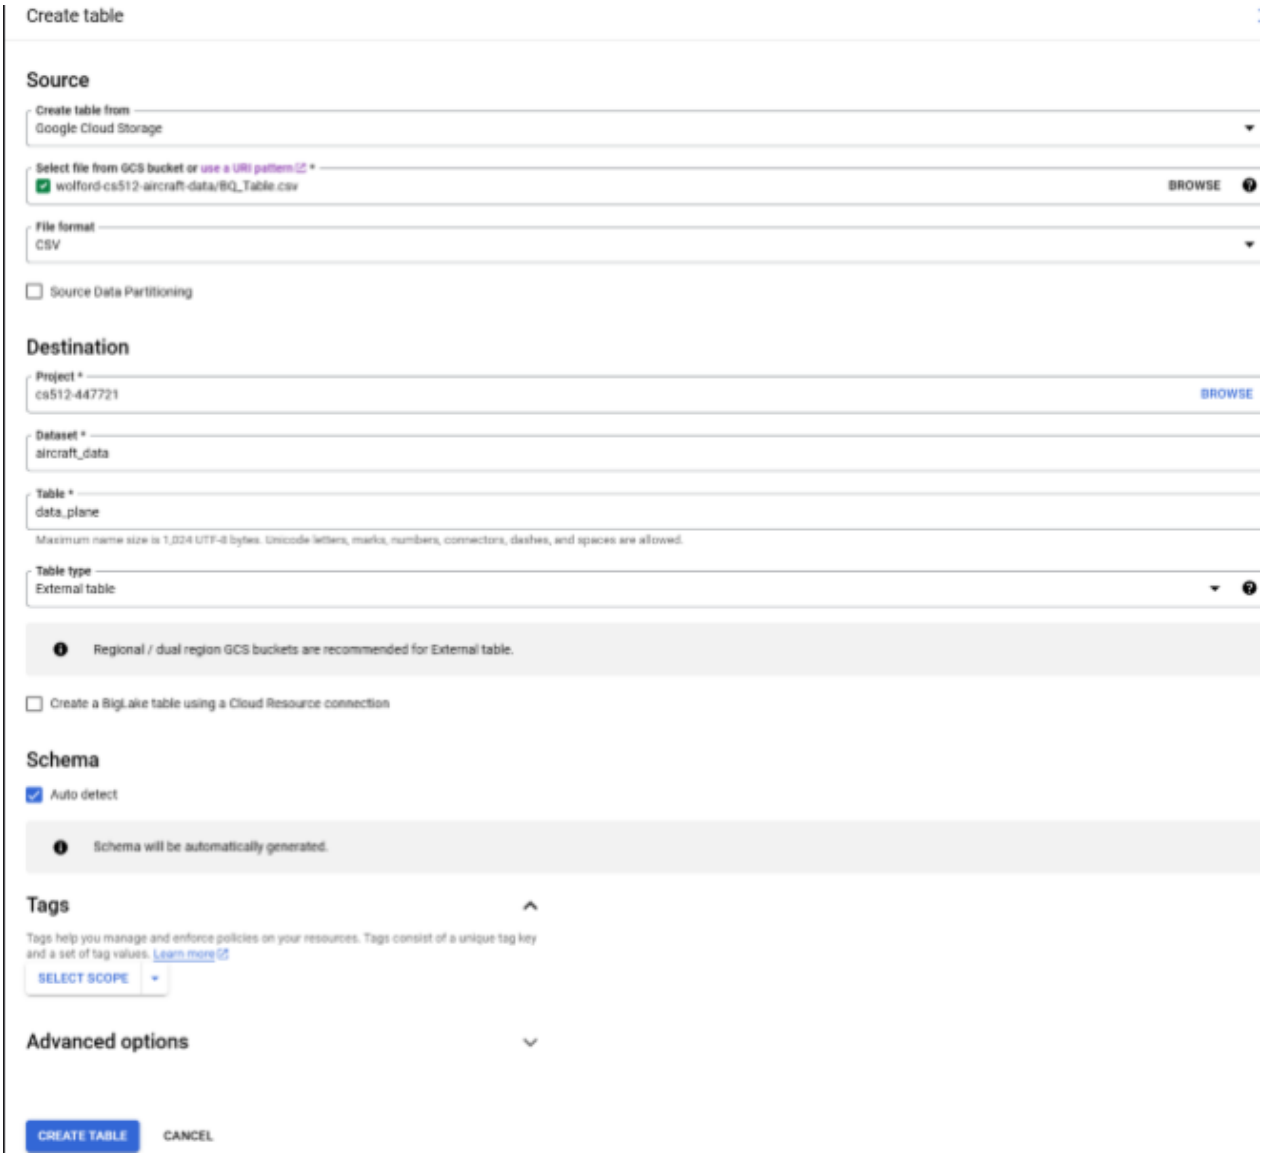
Create table
Source
Create table from
Google Cloud Storage
Select file from GCS bucket or use a URI pattern *
wolford-cs512-aircraft-data/BQ_Table.csv
BROWSE
File format
CSV
Source Data Partitioning
Destination
Project *
cs512-447721
BROWSE
Dataset *
aircraft_data
Table *
data_plane
Maximum name size is 1,024 UTF-8 bytes. Unicode letters, marks, numbers, connectors, dashes, and spaces are allowed.
Table type
External table
Regional / dual region GCS buckets are recommended for External table.
Create a BigLake table using a Cloud Resource connection
Schema
Auto detect
Schema will be automatically generated.
Tags
Tags help you manage and enforce policies on your resources. Tags consist of a unique tag key and a set of tag values. Learn more
SELECT SCOPE
Advanced options
CREATE TABLE
CANCEL

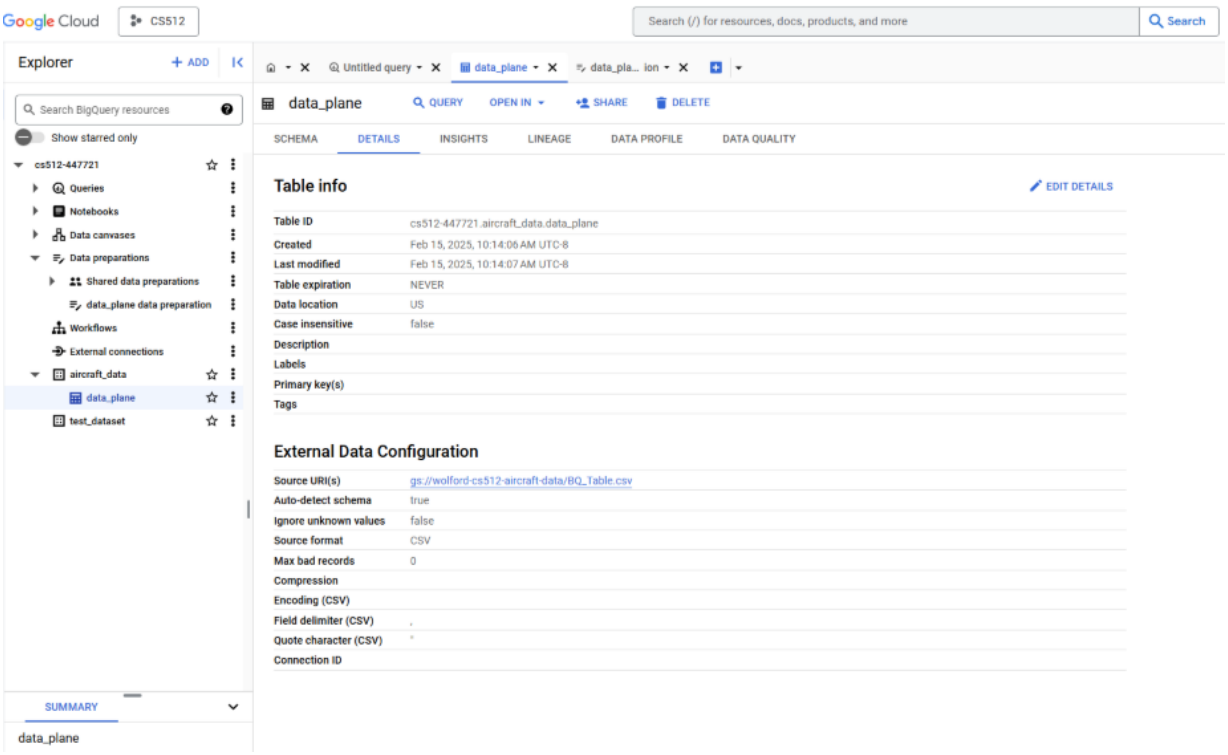

Open in data prep.

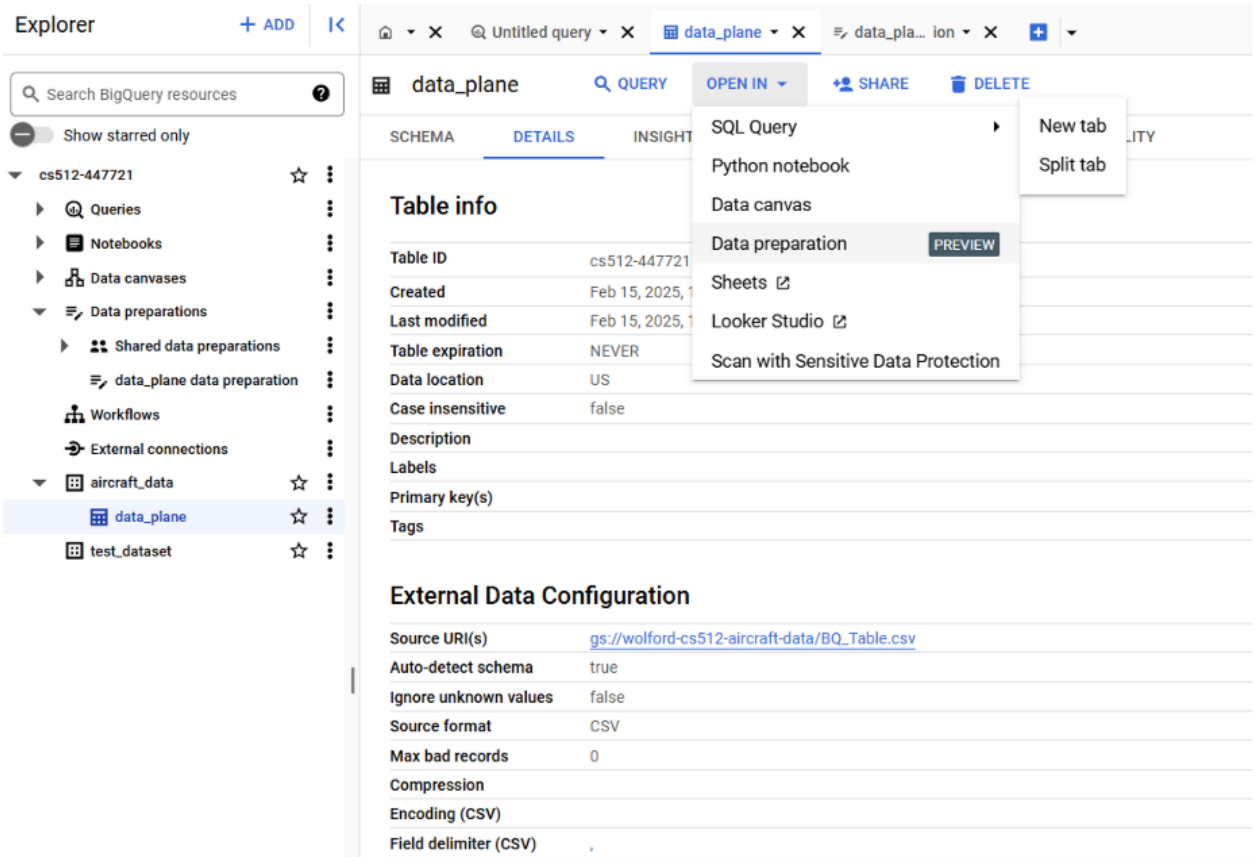

And visualize all 10,000 rows.

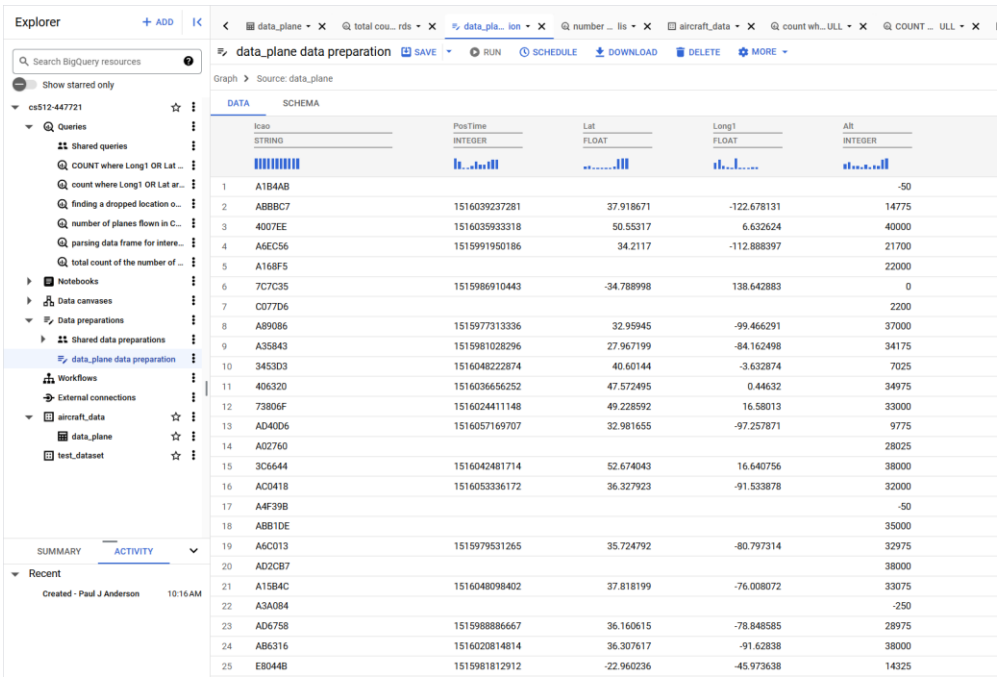

6. Write BigQuery SQL to compute the answer

We then opened a query in Google Big Query and used the given query to answer the question about the data and hit Run.

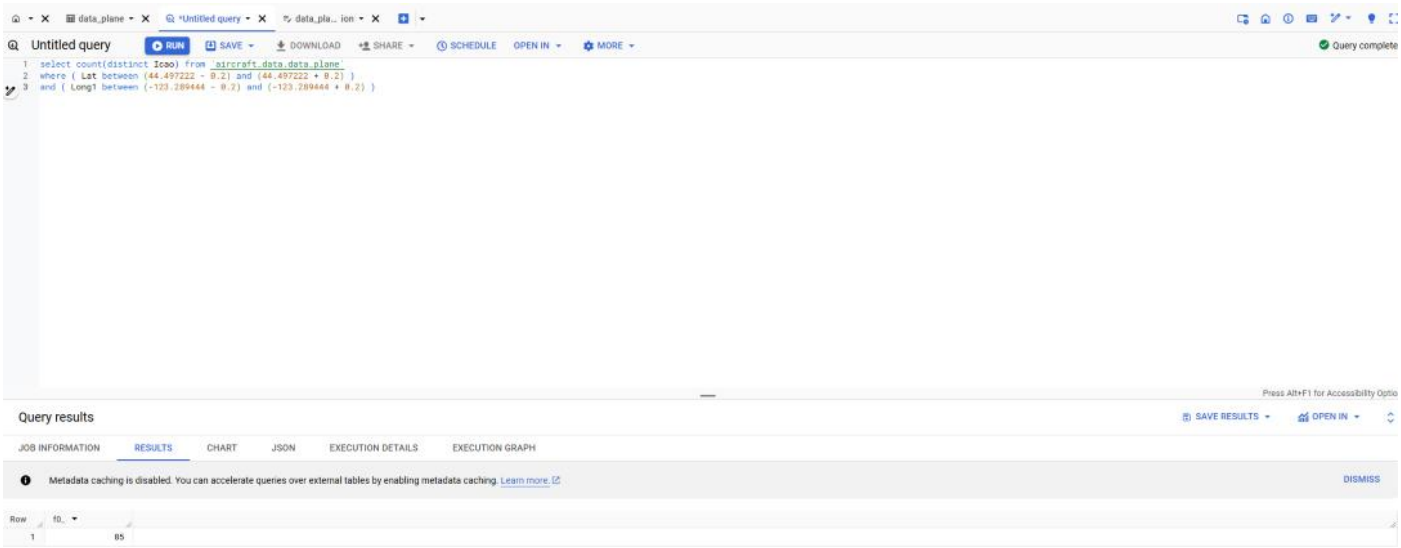

The query returned 85 entries.

I then took some time to explore this large dataset, which seems to be worldwide flights for a set time period. I ran the following queries to explore the data:

How many total rows

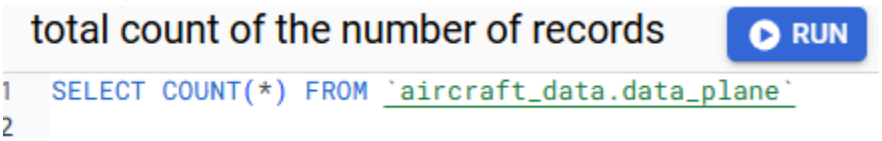

Returns: 14446218

How many cells with empty spaces in Long1 or Lat

count where Long1 OR Lat are NULL RU

```
SELECT COUNT(*) FROM `aircraft_data.data_plane`
WHERE Long1 IS NULL OR Lat IS NULL
```

Returns: 3953109

How many cells with empty spaces in Long1 or Lat or Alt

COUNT where Long1 OR Lat OR Alt... RUN

```
1  SELECT COUNT(*) FROM `aircraft_data.data_plane`
2  WHERE Long1 IS NULL OR Lat IS NULL OR Alt IS NULL
```

Returns: 3991106

The TA then had us run the following Query

parsing data frame for interesting ... RUN SAVE QUERY DOWNLOAD SHARE SCHEDULE OPEN IN MORE This query will proce

```
WITH PlaneData AS ( -- this is defining a function/main query --
    SELECT
        *,
        LAG(Lat) OVER (PARTITION BY Icao ORDER BY PosTime) AS prev_Lat, -- LAG how much distance between previous and current point, previous row --
        LAG(Long1) OVER (PARTITION BY Icao ORDER BY PosTime) AS prev_Long, -- LAG then OVER, PARTITION BY partitions columns and rows, ORDER BY asc desc or PosTime, AS making a variable prev_Lat --
        LAG(PosTime) OVER (PARTITION BY Icao ORDER BY PosTime) AS prev_PosTime
    FROM `aircraft_data.data_plane` -- creates a new data frame from these data --
),
DistanceCalc AS ( -- calc accounts for spherical space and the distance between two points --
    SELECT
        *,
        6371 * 2 * ASIN(
            SQRT(
                POW(SIN((Lat - prev_Lat) * 3.14159 / 360), 2) +
                COS(prev_Lat * 3.14159 / 180) * COS(Lat * 3.14159 / 180) *
                POW(SIN((Long1 - prev_Long) * 3.14159 / 360), 2)
            )
        ) AS distance_km,
        (PosTime - prev_PosTime) AS time_diff
    FROM PlaneData
)
SELECT *
FROM DistanceCalc
WHERE distance_km > 10 -- normalization function used as a filter, or finding the wacky ones --
  AND time_diff < 60;
```

This parsed the colums and rows to create a new data frame with the fields for previous position so that the distance between those two positions can be calculated using the geometry of a sphere. It then sets a filter level for distance (km) and time (seconds).

After running this, we found two interesting patterns.

1. There were many null fields for PosTime, Lat, and Long1
2. There Icao ID A4A0CF seemed to jump in position irregularly.

We ran the following script on the original dataframe to try and find out more about this plane.

finding a dropped location on a uni... RUN

```
1  SELECT * FROM `aircraft_data.data_plane`
2  WHERE Icao =  "A4A0CF"
```

After running this, we found one interesting pattern.

1. Out of the 873 records for this unique Icao, 587 had null records for PosTime, Lat, and Long1
2. 12 of these had null altitude

In considering why this might be the case, we hypothesized that the plane could have had its transmitter turned off or it could have simply malfunctioned midair. With more insight into the instrumentation of aircraft, and transmission of these four variables, with the planes Icao signature, there may be a simple conclusion or more interesting questions to be investigated.

Note: There is some concern that the JSON scrubbing in data prep was not completed. Based on the rubric, this is a significant issue. However, we feel we have documented and communicated about the issue well. Please let our team know if we can learn about this function of Google Cloud Services in another way and we would be happy to review it.

We were both able to complete the assignment in a similar fashion. Paul wrote the report.

Code References

1. Debian. (n.d.). Unix text tools. In *Debian Reference Manual*. Retrieved February 15, 2025, from https://www.debian.org/doc/manuals/debian-reference/ch01.en.html#_unix_text_tools
2. Debian. (n.d.). *FileSystem*. Retrieved February 15, 2025, from https://wiki.debian.org/FileSystem
3. Linux Kernel Labs. (n.d.). *Filesystem*. Retrieved February 15, 2025, from https://linux-kernel-labs.github.io/refs/heads/master/lectures/fs.html
4. Debian. (n.d.). *Debian Reference Manual*. Retrieved February 15, 2025, from https://www.debian.org/doc/manuals/debian-reference/ch01.en.html
5. Debian. (n.d.). *Shell Commands*. Retrieved February 15, 2025, from https://wiki.debian.org/ShellCommands
6. Debian. (n.d.). *Package basics*. In *Debian FAQ*. Retrieved February 15, 2025, from https://www.debian.org/doc/manuals/debian-faq/pkg-basics.en.html
7. Linux.com. (n.d.). *How to use the Linux command line: Basics of CLI*. Retrieved February 15, 2025, from https://www.linux.com/training-tutorials/how-use-linux-command-line-basics-cli/
8. Oregon State University. (n.d.). *Exploration: Data storage walkthrough 2* [Course module]. In CS512: Data Storage. Retrieved February 15, 2025, from https://canvas.oregonstate.edu/courses/1988535/pages/exploration-data-storage-walkthrough-2?module_item_id=25112246